

Task-Scoped Consent for Personal AI Agents:

(greshake2023not, yi2023benchmarking, zhan2024injecagent, debenedetti2024agentdojo,

understandable to non-experts. Non-expert comprehension remains a design target represented

exactly test PAPF's full permission-boundary question. AgentDojo evaluates attacks and

Table 1: Main metrics from paper/tables/main_metrics.md.

artifact, not a full consumer-agent benchmark release.

harder benchmark slices, and compare against faithful implementations or closer